

Audio Systems Learning Map (Local-First, Python-Centric)

This document is a **learning-oriented map of audio projects** worth studying if you want to design **real-time, offline, programmable audio systems**. The emphasis is on **architecture, signal flow, buffers, and extensibility**, not just end-user tools.

Python is prioritized, but strategically chosen C/C++ projects are included when they teach *core engine design* that transfers cleanly to Python bindings.

1. Core Mental Models (Read This First)

These concepts show up in *every serious audio system*, regardless of language or framework. Understanding them deeply will save you months.

1.1 Push vs Pull Audio Graphs

- **Pull-based:** audio driver asks for the next buffer (most common: PortAudio, JACK)
- **Push-based:** producer pushes audio downstream (less common, more fragile)

Why it matters: - Pull-based systems naturally enforce realtime constraints - Push-based systems often drift or glitch under load

Look for this concept in: - JACK client callbacks - JUCE AudioProcessor::processBlock

1.2 Ring / Circular Buffers

Key ideas: - Fixed-size memory - Separate read/write heads - No allocations in realtime paths

Intriguing questions: - How do you handle underruns vs overruns? - Do you drop data or block?

This concept connects audio ↔ MIDI ↔ UI safely.

1.3 Clock Domains

Typical clocks: - Audio sample clock (e.g. 48kHz) - MIDI/event clock (timestamped) - UI / render clock (vsync or timer-driven)

Key insight:

Most bugs are *clock-crossing bugs*, not DSP bugs.

1.4 Event vs Signal Pipelines

- **Signals:** continuous (audio buffers)
- **Events:** discrete (notes, beats, UI actions)

Clean systems never confuse the two.

2. Audio I/O, Buffers, and Realtime Foundations

These projects teach how audio *enters and leaves* your program.

◆ sounddevice (Python)

<https://github.com/spatialaudio/python-sounddevice>

Why it's interesting: - Thin wrapper over PortAudio - Exposes callback-based audio I/O

Concepts to study: - Audio callback lifetimes - Block sizes and latency - Thread priority and dropouts

Intriguing question:

What *cannot* happen inside an audio callback?

◆ JACK Audio (Reference Architecture)

<https://github.com/jackaudio/jack2>

Why it's interesting: - Explicit audio graph scheduling - Sample-accurate synchronization

Concepts to study: - Client registration - Graph execution order - Realtime safety guarantees

Even if you never run JACK, copy its *mental model*.

◆ JACK Audio (Architecture Reference)

(Language: C, bindings everywhere)

What to extract: - Graph-scheduled audio processing - Sample-accurate synchronization - Client registration model

Even if you never use JACK directly, its *mental model* is gold.

3. Music Information Retrieval & Feature Extraction

These projects turn raw audio into *structured data*.

- ◆ **librosa (Python)**

<https://github.com/librosa/librosa>

Why it's interesting: - Reference implementations of DSP

Concepts to focus on: - STFT windowing tradeoffs - Frame vs hop size - Phase reconstruction

Intriguing question:

How much temporal resolution do you really need?

- ◆ **Essentia (C++ core + Python bindings)**

<https://github.com/MTG/essentia>

Why it's interesting: - Descriptor graphs - Streaming vs batch analyzers

Concepts to focus on: - Feature dependency graphs - Deterministic pipelines

Think of this as a *feature engine*, not a library.

- ◆ **madmom (Python)**

<https://github.com/CPJKU/madmom>

Why it's interesting: - Beat/downbeat tracking - Musical-time reasoning

Concepts to focus on: - Tempo hypotheses - Probabilistic timing

- ◆ **Essentia (C++ core + Python bindings)**

Why study: industrial-grade MIR system

Extract ideas: - Descriptor graphs - Feature pipelines - Batch vs streaming analyzers

This is a *blueprint* for turning audio into structured data.

- ◆ **madmom (Python)**

Why study: rhythm-aware systems

Study focus: - Beat/downbeat tracking - Synchronization to musical time - Probabilistic timing models

Excellent for drum trainers and sequencers.

4. Source Separation & Stem Architectures

These projects show how ML integrates with audio I/O.

- ◆ **Ultimate Vocal Remover (UVR)**

<https://github.com/Anjok07/ultimatevocalremovergui>

Why it's interesting: - Chunked inference - Overlap-add reconstruction

Concepts to study: - Window stitching - Model orchestration

Ignore the UI first — read the processing pipeline.

- ◆ **Demucs**

<https://github.com/facebookresearch/demucs>

Why it's interesting: - End-to-end waveform models - Streaming-friendly design

Concepts to study: - Receptive fields - Latency vs quality

- ◆ **Open-Unmix**

<https://github.com/sigsep/open-unmix-pytorch>

Why it's interesting: - Modular ML separation

Concepts to study: - Model abstraction layers - Separation as a service

- ◆ **Open-Unmix**

Why study: modular ML separation design

Extract: - Configurable stems - Model abstraction layers - Separation as a service concept

Good reference for *pluggable* ML audio blocks.

5. Voice, Pitch, and Performance Modeling

These systems convert sound into *control*.

- ◆ **CREPE**

<https://github.com/marl/crepe>

Why it's interesting: - Neural pitch estimation

Concepts to study: - Frame confidence - Audio → control-rate conversion

- ◆ **WORLD Vocoder**

<https://github.com/mmorise/World>

Why it's interesting: - Clean speech decomposition

Concepts to study: - Separation of pitch, timbre, noise - Deterministic synthesis

This explains how voice cloning actually works.

- ◆ **DDSP**

<https://github.com/magenta/ddsp>

Why it's interesting: - Neural networks parameterize DSP

Concepts to study: - Hybrid systems - Controllability vs realism

- ◆ **WORLD Vocoder**

Why study: clean speech decomposition

Extract: - Separation of pitch, timbre, noise - Deterministic synthesis

This teaches how voice cloning systems *actually work under the hood.*

- ◆ **DDSP**

Why study: hybrid neural + DSP systems

Key insight: - Neural networks don't replace DSP — they *parameterize it*

Highly relevant if you care about controllability.

6. MIDI, Timing, and Event Systems

MIDI is not audio — treat it as events.

- ◆ **RtMidi**

<https://github.com/thestk/rtmidi>

Why it's interesting: - Cross-platform MIDI I/O

Concepts to study: - Polling vs callbacks - Timestamping

- ◆ **mido**

<https://github.com/mido/mido>

Why it's interesting: - Clean Python MIDI abstraction

Concepts to study: - Backend separation - Message routing

- ◆ **pretty_midi**

<https://github.com/craffel/pretty-midi>

Why it's interesting: - MIDI ↔ time alignment

Concepts to study: - Symbolic vs real-time representations

- ◆ **pretty_midi**

Why study: MIDI ↔ time alignment

Extract: - MIDI as structured score data - Mapping note events to wall-clock time

Excellent bridge between symbolic and signal domains.

7. Realtime DSP Graphs (Non-Python, Worth Studying)

- ◆ **SuperCollider**

Why it matters: - Audio as message-passing - Explicit separation of control vs audio rate

This will change how you design Python engines.

- ◆ **Faust**

Why it matters: - Pure functional DSP - Explicit signal graphs - Compile-time optimization

Faust teaches *what DSP code should look like before glue exists.*

- ◆ **JUCE**

Why it matters: - Production-grade plugin architecture - Audio threading discipline

Study it for architecture, not syntax.

8. Visualization & Audio-UI Coupling

Recommended pairing ideas: - PyQtGraph + ring buffers - OpenGL textures fed from audio FFTs - Timeline-based waveform caches

Study goal: - Decouple rendering from audio time - UI reads snapshots, never live buffers

9. Suggested Learning Path (Practical)

Phase 1 — Fundamentals - sounddevice + ring buffer - librosa offline analysis

Phase 2 — Events & Timing - MIDI input + timestamps - madmom beat tracking

Phase 3 — ML Blocks - Demucs or UVR as a batch service - CREPE pitch → MIDI

Phase 4 — System Design - Audio + MIDI unified event bus - Feature extractors as producers - UI panels as observers

10. What to Build While Learning

Small but powerful projects:

- Realtime pitch → MIDI monitor
- Stem separation → per-stem FFT viewer
- Drum hit alignment trainer
- Audio feature recorder → Zarr store

Each one reinforces architecture, not just DSP.

11. Design Rule of Thumb

If audio can block, it must not allocate. If UI can stall, it must never touch live audio. If ML is slow, treat it as an offline or chunked service.

12. Where This Naturally Leads

This learning map naturally supports: - Modular audio engines - AI-assisted music tools - Drum and vocal training systems - Audio-driven visualization engines - Audio-LLM hybrid pipelines

If you want next: - A **repo reading order** for any section - A **minimal Python audio engine skeleton** - A **unified audio + MIDI event model**

Say the word and we'll go there.