

Procedural UI Rendering & Animation in ModernGL

A Style-Token + NodeGraph + GPU Instance Buffer Architecture

Date: 2026-01-17

Abstract

This paper describes a practical, GPU-friendly architecture for building a stylized, animated UI like the provided dial-and-wave mockup using ModernGL. Instead of redrawing or editing image pixels each frame, the system treats UI elements as data-driven nodes whose appearance is defined by a small set of style tokens and whose motion is driven by animation tracks. Rendering is performed with instanced quads into an offscreen framebuffer (FBO), while fragment shaders procedurally draw rings, quarter-arcs, glows, and waveform strokes. The result is a decoupled, high-performance UI layer that can be composed into a larger engine, cached when idle, and re-themed instantly by swapping a single token object.

1. Design Goals

Goals: (1) preserve the mockup's look via reusable style tokens, (2) animate by changing parameters rather than pixels, (3) decouple UI layout/data/animation from rendering, (4) render to a texture buffer for compositing, (5) make it extensible and Python-idiomatic.

Non-goals: a full widget toolkit; full text shaping; perfect font atlas management; arbitrary vector graphics. The focus is the core engine pattern: tokens → nodes → tracks → GPU buffers → procedural shaders → FBO.

2. Core Idea: UI as Instanced Quads + Procedural Fragment Shaders

Every UI component (dial, wave lane, histogram panel) is drawn as a screen-aligned quad. A per-instance parameter block describes the component rectangle and its visual state. The fragment shader then renders the component procedurally inside that rectangle using analytic geometry and smoothstep-based anti-aliasing.

This avoids per-frame texture edits. Instead, you upload a small instance buffer every frame (or only when dirty) and let the GPU synthesize the final pixels.

2.1 Render-to-Texture UI Buffer

Render the entire UI into an offscreen framebuffer (FBO) with a color texture attachment (`ui_tex`). The engine can then composite `ui_tex` onto a 3D surface, overlay it in the final pass, or feed it to a panel compositor.

3. Data Model: Style Tokens, Nodes, and Properties

3.1 Style Tokens (Design System)

Style tokens are a compact set of semantic values: background colors, stroke alphas, accent hues, glow strengths, radii, spacing, and typography scaling. Tokens are not component-specific; they define the

global style vocabulary. The GPU reads tokens from a global uniform buffer (UBO) called StyleUBO.

```
# tokens.py
from dataclasses import dataclass
from typing import Tuple

Vec4 = Tuple[float, float, float, float]

@dataclass(frozen=True)
class StyleTokens:
    bg_app: Vec4
    panel: Vec4
    well: Vec4
    stroke1: Vec4
    stroke2: Vec4
    text1: Vec4
    text2: Vec4
    accent_teal: Vec4
    accent_gold: Vec4
    accent_green: Vec4
    accent_peri: Vec4
    glow_core: float
    glow_halo: float

DEFAULT_STYLE = StyleTokens(
    bg_app=(0.14, 0.15, 0.16, 1.0),
    panel=(0.17, 0.18, 0.20, 1.0),
    well=(0.11, 0.12, 0.14, 1.0),
    stroke1=(1.00, 1.00, 1.00, 0.10),
    stroke2=(1.00, 1.00, 1.00, 0.06),
    text1=(0.85, 0.84, 0.82, 1.0),
    text2=(0.60, 0.62, 0.65, 1.0),
    accent_teal=(0.68, 0.91, 0.87, 1.0),
    accent_gold=(0.89, 0.70, 0.42, 1.0),
    accent_green=(0.58, 0.69, 0.51, 1.0),
    accent_peri=(0.63, 0.65, 0.78, 1.0),
    glow_core=0.55,
    glow_halo=0.18,
)
```

3.2 Nodes (UI Scene Graph as Data)

A UI node is a plain data object: kind, rectangle, instance id, and a typed property bag. Nodes do not render themselves; renderers consume nodes and produce GPU instance data. This keeps the model decoupled from ModernGL and allows worker threads to update node properties without touching the GL context.

```
# nodes.py
from dataclasses import dataclass, field
from typing import Dict, Any, Tuple

Rect = Tuple[float, float, float, float] # x, y, w, h in pixels

@dataclass
class UINode:
    kind: str
    rect: Rect
    instance_id: int
    props: Dict[str, Any] = field(default_factory=dict)
    def getf(self, key: str, default: float = 0.0) -> float:
```

```

        v = self.props.get(key, default)
        return float(v)

    def set(self, key: str, value: Any) -> None:
        self.props[key] = value

```

4. Animation: Tracks Compiled to Fast Setters

Animation is stored as tracks targeting node properties by path. The key trick is compiling paths into closures once at load time. At runtime, evaluating a track becomes: compute value → call setter. This avoids string parsing every frame.

```

# animation.py
from dataclasses import dataclass
from typing import Callable, List, Tuple, Any
import bisect
import math

Key = Tuple[float, Any]

@dataclass
class Track:
    length: float
    keys: List[Key]
    setter: Callable[[Any], None]
    interp: str = "linear"

    def eval(self, t: float) -> None:
        if not self.keys:
            return
        if t <= self.keys[0][0]:
            self.setter(self.keys[0][1]); return
        if t >= self.keys[-1][0]:
            self.setter(self.keys[-1][1]); return

        times = [k[0] for k in self.keys]
        i = bisect.bisect_right(times, t) - 1
        t0, v0 = self.keys[i]
        t1, v1 = self.keys[i + 1]
        u = (t - t0) / max(1e-9, (t1 - t0))

        if self.interp == "step":
            self.setter(v0); return

        # numeric-only linear interpolation
        if isinstance(v0, (int, float)) and isinstance(v1, (int, float)):
            self.setter((1.0 - u) * float(v0) + u * float(v1)); return

        # vec4 interpolation
        if isinstance(v0, (list, tuple)) and len(v0) == 4:
            out = [(1.0 - u)*v0[j] + u*v1[j] for j in range(4)]
            self.setter(out); return

        self.setter(v0)

@dataclass
class Clip:
    name: str
    duration: float

```

```

tracks: List[Track]

def apply(self, t: float) -> None:
    t = t % self.duration
    for tr in self.tracks:
        tr.eval(t)

def compile_path(nodes_by_name: dict, path: str) -> Callable[[Any], None]:
    # Example paths: "dial0.needle", "wave.opacity"
    node_name, prop = path.split(".", 1)
    node = nodes_by_name[node_name]
    return lambda v: node.set(prop, v)

```

5. GPU Data: Instance Struct and Packing

The rendering core is an instance buffer containing one record per component. A simple layout uses five vec4 values (aligned, cache-friendly): rect, accent, stroke, misc0, misc1. Rect is pixel-space (x,y,w,h). Misc fields hold kind-specific parameters such as needle angle and arc lengths.

```

# packing.py
import numpy as np

# Instance layout: 5x vec4 floats = 20 floats per instance
# [0] rect:  x y w h
# [1] accent: r g b a
# [2] stroke: r g b a
# [3] misc0:  needle_angle, opacity, glow, reserved
# [4] misc1:  arc0, arc1, arc2, arc3

STRIDE_FLOATS = 20

def pack_dial(node, style):
    x,y,w,h = node.rect
    rect = (x,y,w,h)

    accent = node.props.get("accent", style.accent_teal)
    stroke = style.stroke1

    needle = node.getf("needle", 0.0)
    opacity = node.getf("opacity", 1.0)
    glow = node.getf("glow", style.glow_core)

    arcs = node.props.get("arcs", (0.25,0.25,0.25,0.25))

    misc0 = (needle, opacity, glow, 0.0)
    misc1 = arcs
    return rect, accent, stroke, misc0, misc1

def pack_wave(node, style):
    x,y,w,h = node.rect
    rect = (x,y,w,h)

    accent = node.props.get("accent", style.accent_peri)
    stroke = style.stroke2

    amp = node.getf("amp", 0.35)
    phase = node.getf("phase", 0.0)
    opacity = node.getf("opacity", 1.0)
    misc0 = (amp, phase, opacity, 0.0)

```

```

    misc1 = (0.0,0.0,0.0,0.0)
    return rect, accent, stroke, misc0, misc1

PACKERS = {
    "dial": pack_dial,
    "wave": pack_wave,
}

def pack_instances(nodes, style):
    out = np.zeros((len(nodes), STRIDE_FLOATS), dtype=np.float32)
    for i, node in enumerate(nodes):
        rect, accent, stroke, misc0, misc1 = PACKERS[node.kind](node, style)
        out[i, 0:4] = rect
        out[i, 4:8] = accent
        out[i, 8:12] = stroke
        out[i, 12:16] = misc0
        out[i, 16:20] = misc1
    return out.reshape(-1)

```

6. Rendering: Instanced Quad + Procedural Shaders

Renderers draw a unit quad per instance. Vertex shader expands the quad into the instance rectangle in clip space. Fragment shader computes local coordinates and draws rings, arcs, crosshair lines, dots, and glows with smoothstep antialiasing. Separate programs per kind (dial, wave) keep shaders simple and extensible.

6.1 Instanced Quad Vertex Shader (Shared)

```

// ui_quad.vert
#version 330

in vec2 in_pos;           // [-1..+1] quad corners or [0..1] depending on convention
in vec2 in_uv;

uniform vec2 u_resolution;

layout(std430, binding = 0) buffer Instances {
    vec4 rects[]; // for demo: pack rect at vec4[instance*5 + 0]
};

out vec2 v_uv;
flat out int v_instance;

void main() {
    v_uv = in_uv;
    v_instance = gl_InstanceID;

    vec4 r = rects[v_instance * 5 + 0]; // x y w h in pixels
    vec2 px = vec2(r.x, r.y) + in_uv * vec2(r.z, r.w); // local->pixel
    vec2 ndc = (px / u_resolution) * 2.0 - 1.0;
    ndc.y *= -1.0; // if your pixel origin is top-left

    gl_Position = vec4(ndc, 0.0, 1.0);
}

```

6.2 Dial Fragment Shader (Quarter Ring + Needle + Glow)

```

// dial.frag
#version 330

uniform vec2 u_resolution;

layout(std430, binding = 0) buffer Instances {
    vec4 data[]; // 5 vec4 per instance
};

in vec2 v_uv;
flat in int v_instance;

out vec4 fragColor;

float aastep(float edge0, float edge1, float x) {
    return smoothstep(edge0, edge1, x);
}

void main() {
    int base = v_instance * 5;
    vec4 rect    = data[base + 0];
    vec4 accent  = data[base + 1];
    vec4 stroke  = data[base + 2];
    vec4 misc0   = data[base + 3];
    vec4 arcs    = data[base + 4];

    float needle = misc0.x;
    float opacity = misc0.y;
    float glow = misc0.z;

    // local coord in [-1..1] with square aspect
    vec2 p = (v_uv * 2.0 - 1.0);
    float r = length(p);
    float ang = atan(p.y, p.x); // [-pi..pi]
    if (ang < 0.0) ang += 6.28318530718;

    // ring parameters
    float ringR = 0.78;
    float ringT = 0.06;
    float ring = aastep(ringR + ringT, ringR, r) * aastep(ringR - ringT, ringR, r);

    // quarter arcs: four spans each of pi/2, with per-quarter length scaling
    float q = 1.57079632679; // pi/2
    float arcMask = 0.0;
    for (int k=0; k<4; k++) {
        float a0 = float(k) * q;
        float a1 = a0 + q * arcs[k];
        float inA = step(a0, ang) * step(ang, a1);
        arcMask = max(arcMask, inA);
    }

    // crosshair lines
    float vx = 1.0 - smoothstep(0.0, 0.01, abs(p.x));
    float vy = 1.0 - smoothstep(0.0, 0.01, abs(p.y));
    float cross = max(vx, vy) * 0.25;

    // needle: distance to a ray
    vec2 dir = vec2(cos(needle), sin(needle));
    float proj = dot(p, dir);
    vec2 perp = p - dir * proj;
}

```

```

float d = length(perp);
float needleMask = smoothstep(0.02, 0.0, d) * smoothstep(-0.2, 0.4, proj);

// base color
vec3 col = vec3(0.10);
col += ring * stroke.rgb * 0.65;
col += ring * arcMask * accent.rgb * 1.25;
col += needleMask * accent.rgb * 1.10;
col += cross * stroke.rgb;

// glow (cheap)
float glowRing = smoothstep(ringR + ringT + 0.06, ringR + ringT, r) * ring;
col += glowRing * accent.rgb * glow;

fragColor = vec4(col, opacity);
}

```

6.3 Wave Fragment Shader (Line + Glow)

```

// wave.frag
#version 330

layout(std430, binding = 0) buffer Instances {
    vec4 data[]; // 5 vec4 per instance
};

in vec2 v_uv;
flat in int v_instance;
out vec4 fragColor;

float waveFunc(float x, float amp, float phase) {
    // demo wave; replace with SSBO-sampled data for real waveform
    return 0.5 + amp * sin(6.28318 * x + phase);
}

void main() {
    int base = v_instance * 5;
    vec4 accent = data[base + 1];
    vec4 stroke = data[base + 2];
    vec4 misc0 = data[base + 3];

    float amp = misc0.x;
    float phase = misc0.y;
    float opacity = misc0.z;

    float y = waveFunc(v_uv.x, amp, phase);
    float d = abs(v_uv.y - y);

    // line thickness
    float line = smoothstep(0.015, 0.0, d);
    float glow = smoothstep(0.06, 0.0, d) * 0.35;

    vec3 col = vec3(0.06);
    col += line * accent.rgb * 1.1;
    col += glow * accent.rgb * 0.6;

    fragColor = vec4(col, opacity);
}

```

7. Implementation Skeleton: Putting It Together

This section provides a minimal runnable structure. It assumes moderngl_window for context creation. The code demonstrates: style UBO update, instance buffer packing, dial + wave render, and a simple animation that drives needle and wave phase.

```
# demo_app.py (skeleton)
import math
import numpy as np
import moderngl
import moderngl_window as mglw

from tokens import DEFAULT_STYLE
from nodes import UINode
from packing import pack_instances, STRIDE_FLOATS

class Demo(mglw.WindowConfig):
    gl_version = (3, 3)
    title = "Procedural UI Dial + Wave (ModernGL)"
    window_size = (1100, 700)
    aspect_ratio = None
    resizable = True

    def __init__(self, **kwargs):
        super().__init__(**kwargs)
        self.style = DEFAULT_STYLE

        # Build nodes
        self.nodes = []
        self.nodes.append(UINode("dial", (60, 60, 220, 220), 0, {
            "needle": 0.2, "arcs": (0.25, 0.25, 0.25, 0.25), "accent": self.style.accent_teal
        }))
        self.nodes.append(UINode("dial", (310, 60, 220, 220), 1, {
            "needle": 1.0, "arcs": (0.40, 0.15, 0.25, 0.30), "accent": self.style.accent_gold
        }))
        self.nodes.append(UINode("wave", (60, 330, 800, 140), 2, {
            "amp": 0.22, "phase": 0.0, "opacity": 1.0, "accent": self.style.accent_peri
        }))

        # Compile programs
        self.quad_vbo = self.ctx.buffer(np.array([
            # pos      uv
            0.0, 0.0, 0.0, 0.0,
            1.0, 0.0, 1.0, 0.0,
            1.0, 1.0, 1.0, 1.0,
            0.0, 0.0, 0.0, 0.0,
            1.0, 1.0, 1.0, 1.0,
            0.0, 1.0, 0.0, 1.0,
        ], dtype="f4").tobytes())

        self.dial_prog = self.ctx.program(
            vertex_shader=open("ui_quad.vert", "r", encoding="utf8").read(),
            fragment_shader=open("dial.frag", "r", encoding="utf8").read()
        )
        self.wave_prog = self.ctx.program(
            vertex_shader=open("ui_quad.vert", "r", encoding="utf8").read(),
            fragment_shader=open("wave.frag", "r", encoding="utf8").read()
        )

    # Instance buffer: 5 vec4 per instance => 20 floats per instance
```

```

        self.instance_buf = self.ctx.buffer(reserve=len(self.nodes) * STRIDE_FLOATS * 4)

        # Separate VAOs (same VBO + instance SSBO binding)
        self.dial_vao = self.ctx.vertex_array(
            self.dial_prog,
            [(self.quad_vbo, "2f 2f", "in_pos", "in_uv")]
        )
        self.wave_vao = self.ctx.vertex_array(
            self.wave_prog,
            [(self.quad_vbo, "2f 2f", "in_pos", "in_uv")]
        )

    self.t = 0.0

    def render(self, time: float, frame_time: float):
        self.t += frame_time

        # Animate needle + wave phase
        self.nodes[0].set("needle", 0.4 + 0.6 * math.sin(self.t * 1.2))
        self.nodes[1].set("needle", 1.4 + 0.5 * math.sin(self.t * 0.7))
        self.nodes[2].set("phase", self.t * 2.0)

        # Pack + upload
        packed = pack_instances(self.nodes, self.style)
        self.instance_buf.write(packed.tobytes())

        # Bind SSBO to binding=0 for programs
        self.instance_buf.bind_to_storage_buffer(binding=0)

        w, h = self wnd.size
        self.ctx.viewport = (0, 0, w, h)
        self.ctx.clear(0.08, 0.09, 0.10, 1.0)

        # Set resolution uniforms
        self.dial_prog["u_resolution"].value = (w, h)
        self.wave_prog["u_resolution"].value = (w, h)

        # Render: dials are instances 0..1, wave is instance 2
        # Easiest: render all and let shaders interpret their slots; production would batch by kind
        self.dial_vao.render(instances=2)
        self.wave_vao.render(instances=1)

    if __name__ == "__main__":
        mglw.run_window_config(Demo)

```

8. Extensibility: New Components, Batching, and Caching

To add a component, define a packer function and a fragment shader. For performance, batch nodes by kind so each renderer draws a contiguous range of instances. For efficiency, track a global UI dirty flag: if no animations or data changed, skip the UI pass and reuse last ui_tex.

9. Threading: Safe Worker Updates

Workers may update data models (timing offsets, histogram bins, transcription events) without touching OpenGL. They publish results to the main thread via a lock-free queue or a message bus. The main thread applies updates to node properties and then renders. This matches OpenGL's single-context thread constraints.

10. Summary Checklist

A minimal production-ready pipeline looks like:

- StyleTokens → StyleUBO (theme switch in one write)
- Nodes (kind + rect + props) as decoupled data
- Anim tracks compiled to setters (no runtime string parsing)
- Pack nodes into an instance SSBO (std430, contiguous)
- Instanced quad renderer per component kind (dial, wave, histogram)
- Offscreen FBO output (ui_tex) composited into engine
- Dirty caching (skip redraw when idle)

This architecture produces the mockup-like look by relying on consistent style tokens and procedural drawing. It remains Python-idiomatic (dataclasses, closures, dictionaries of callables) while aligning with GPU best practices (instancing, SSBOs, small per-frame uploads, render-to-texture compositing).